

# Problem 1

September 11, 2018

## 0.0.1 Author: Yasir Abdurrahman

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

In [2]: df_train = pd.read_csv('dataset1/train1.csv')
```

## 0.1 Dataset

### 0.1.1 Description

Dataset berikut merupakan data demografi 1556 nasabah bank pada kuartier ke-4 tahun 2017, data ini kemudian dibagi menjadi data train dan test. Variabel target pada data ini adalah `not_paid`, variabel biner yang menjadi indikasi suatu loan/ pinjaman lunas (berhasil dilunasi pembayarannya) atau tidak. Suatu pinjaman dikatakan `not_paid` (`not paid = 1`) jika terjadi default (gagal bayar), Charged Off, atau lewat batas akhir pembayaran (Grace Period).

### 0.1.2 Variables Glossary

- `initial_list_status`: indikasi loan termasuk ke dalam kategori `w` (whole) atau `f` (fractional).
- `purpose`: tujuan peminjaman (loan) terbagi atas 5 kategori yaitu untuk `credit_card`, `debt_consolidation`, `home_improvement`, `major_purchase`, dan `small_bussiness`
- `int_rate`: interest rate(suku bunga) dalam prosentase
- `installment`: banyaknya installment/uang bulanan yang dibayarkan peminjam
- `annual_inc`: income/pemasukan tahunan peminjam sesuai yang tertulis saat proses pengajuan pinjaman
- `dti`: rasio antara pinjaman bulanan yang wajib dibayarkan peminjam dengan gaji/pemasukan peminjam sesuai report
- `verification_status`: status verifikasi report pemasukan/gaji peminjam, terbagi atas kategori `income_verified`, `not verified`, atau `source was verified`
- `grade`: grade loan berdasarkan software
- `revol_bal`: total kredit dalam revolving balance (pinjaman yang tidak terbayarkan)
- `inq_last_12m`: banyaknya kredit/pinjaman pada akhir bulan 12
- `delinq_2yrs`: banyaknya hari telat bayar untuk kriteria 30+ pada history peminjam selama 2 tahun terakhir
- `home_ownership`: kategori kepemilikan rumah peminjam meliputi `MORTGAGE`, `OWN`, atau `RENT`
- `log_inc`: log dari `annual_inc`

- verified: 0 untuk not\_verified masih dibawah status verifikasi, 1 lainnya
- grdCtoA: 1 untuk grade kredit A, B atau C; 0 untuk grade load lainnya
- not\_paid: 1 jika gagal bayar *charge off/grace period*, 0 lainnya (TARGET)

In [3]: df\_train.head()

```
Out[3]:
```

	Unnamed: 0	initial_list_status		purpose	int_rate	installment	\
0	1495	w		debt_consolidation	21.45	955.75	
1	266	w		debt_consolidation	18.06	289.47	
2	309	w		home_improvement	9.44	838.91	
3	239	w		home_improvement	10.42	214.55	
4	136	f		debt_consolidation	11.99	1024.52	

	annual_inc	dti	verification_status	grade	revol_bal	inq_last_12m	\
0	90000.0	20.91	Verified	D	23448	4	
1	65000.0	12.74	Source Verified	D	13362	2	
2	97400.0	12.64	Source Verified	B	2372	1	
3	60000.0	2.38	Not Verified	B	4705	2	
4	150000.0	20.84	Not Verified	B	14342	0	

	delinq_2yrs	home_ownership	not_paid	log_inc	verified	grdCtoA
0	0	MORTGAGE	1	11.407565	1	0
1	0	MORTGAGE	1	11.082143	1	0
2	0	MORTGAGE	1	11.486581	1	1
3	0	OWN	0	11.002100	0	1
4	0	MORTGAGE	0	11.918391	0	1

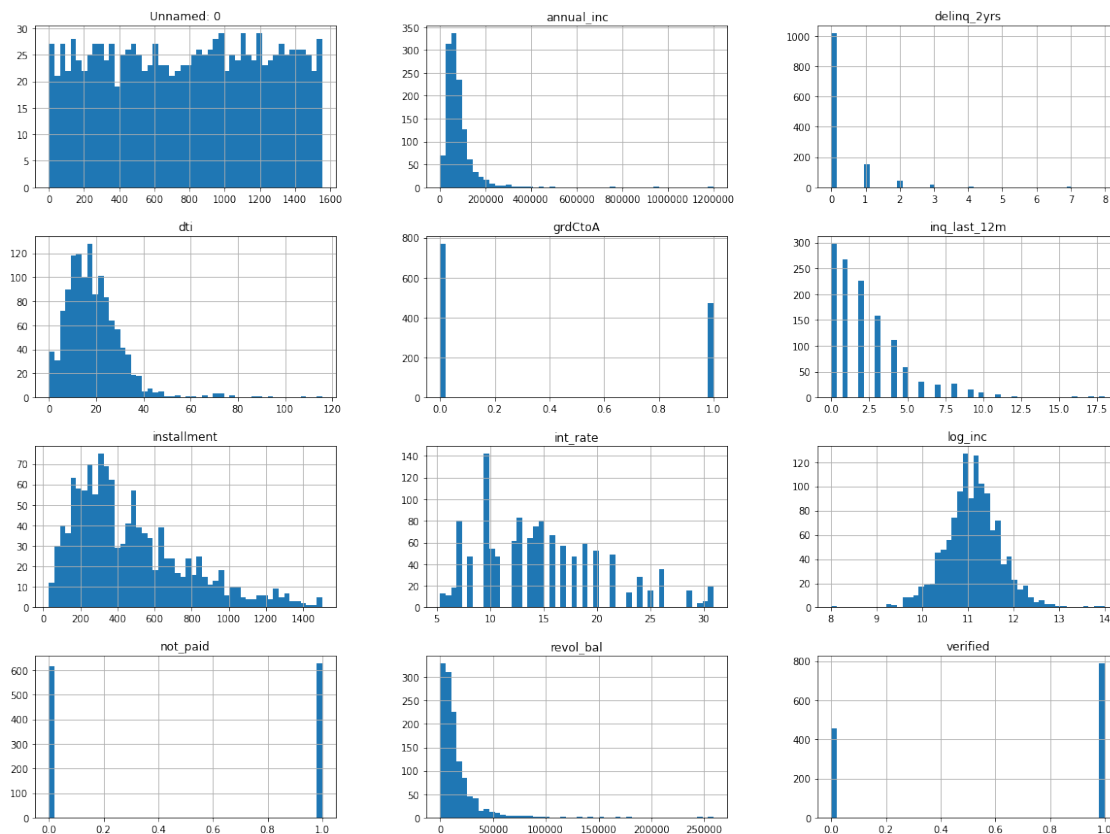
## 1 Explore Dataset

In [4]: df\_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1244 entries, 0 to 1243
Data columns (total 17 columns):
Unnamed: 0      1244 non-null int64
initial_list_status  1244 non-null object
purpose         1244 non-null object
int_rate        1244 non-null float64
installment     1244 non-null float64
annual_inc      1244 non-null float64
dti             1244 non-null float64
verification_status  1244 non-null object
grade          1244 non-null object
revol_bal       1244 non-null int64
inq_last_12m    1244 non-null int64
delinq_2yrs     1244 non-null int64
home_ownership  1244 non-null object
not_paid        1244 non-null int64
```

```
log_inc          1244 non-null float64
verified         1244 non-null int64
grdCtoA         1244 non-null int64
dtypes: float64(5), int64(7), object(5)
memory usage: 165.3+ KB
```

```
In [5]: %matplotlib inline
df_train.hist(bins=50, figsize=(20,15))
plt.show()
```



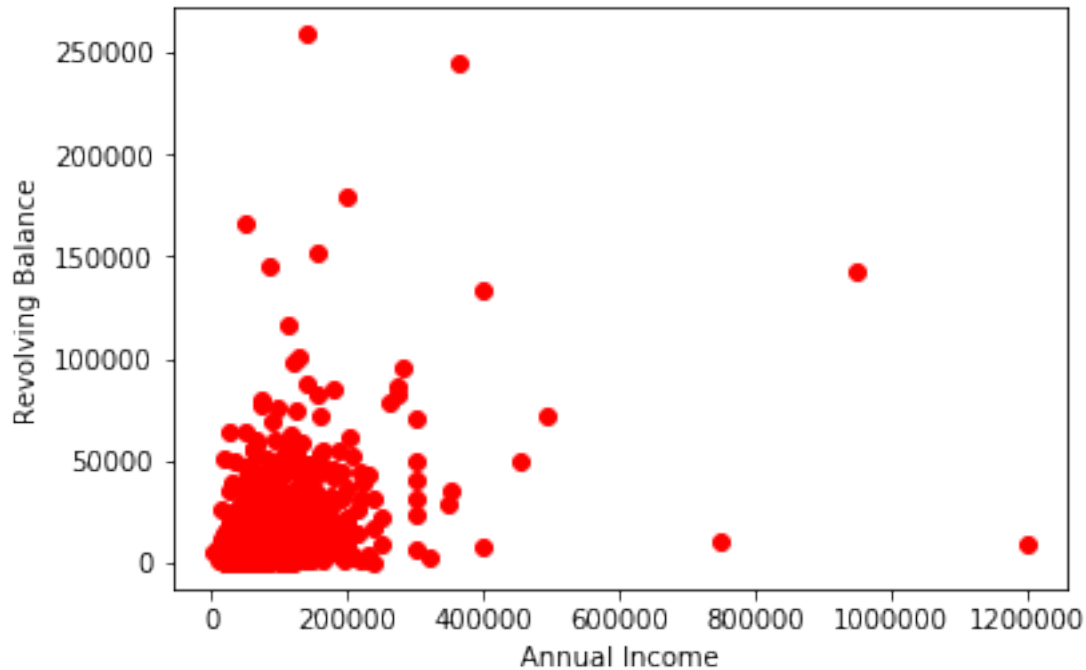
## 1.1 Question 1:

Bagaimana anda mendeskripsikan hubungan antara annual\_inc dan revol\_bal?

**Answer:** Tidak ada hubungan linier

```
In [6]: # plot 'annual_inc' dan 'revol_bal'
plt.plot(df_train['annual_inc'], df_train['revol_bal'], 'ro')
plt.xlabel('Annual Income')
plt.ylabel('Revolving Balance')
```

```
Out[6]: Text(0,0.5,'Revolving Balance')
```



## 1.2 Question 2:

Berdasarkan kategori purpose (tujuan pinjaman) yang paling banyak ditemukan nasabah mengalami gagal bayar(not\_paid=1), berapa banyak nasabah yang mengalami gagal bayar pada kategori tersebut?

**Answer:** debt\_consolidation -> 420

```
In [7]: purpose_not_paid = []

        for index, row in df_train.iterrows():
            if row['not_paid'] == 1:
                purpose_not_paid.append(row['purpose'])
```

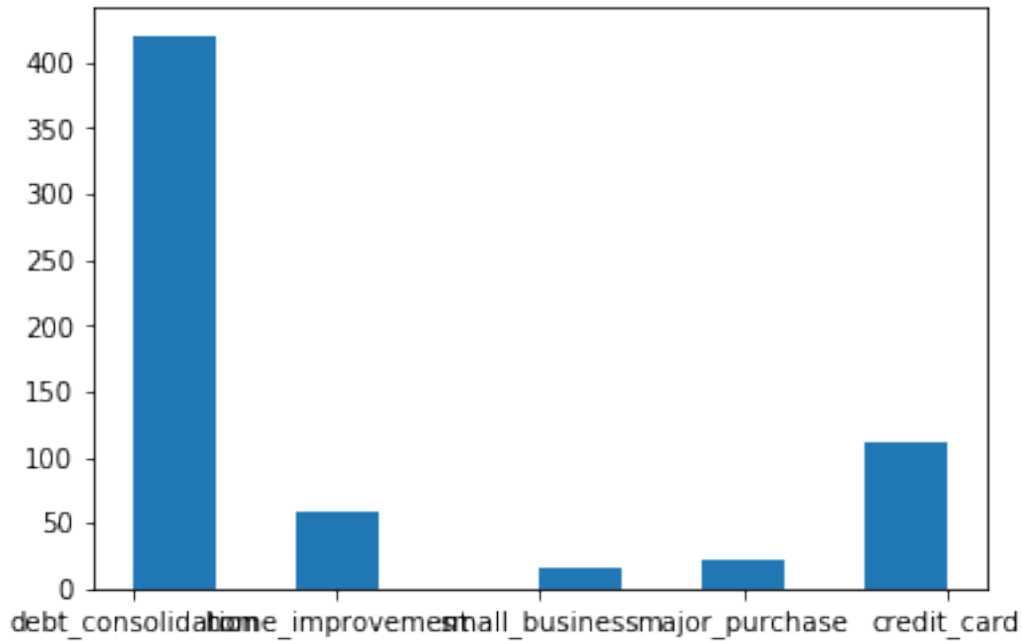
```
In [8]: from collections import Counter

        Counter(purpose_not_paid)
```

```
Out[8]: Counter({'debt_consolidation': 420,
                 'home_improvement': 58,
                 'small_business': 15,
                 'major_purchase': 22,
                 'credit_card': 112})
```

```
In [9]: plt.hist(purpose_not_paid)
```

```
Out[9]: (array([420.,  0.,  58.,  0.,  0.,  15.,  0.,  22.,  0., 112.]),
         array([0. , 0.4, 0.8, 1.2, 1.6, 2. , 2.4, 2.8, 3.2, 3.6, 4. ]),
         <a list of 10 Patch objects>)
```



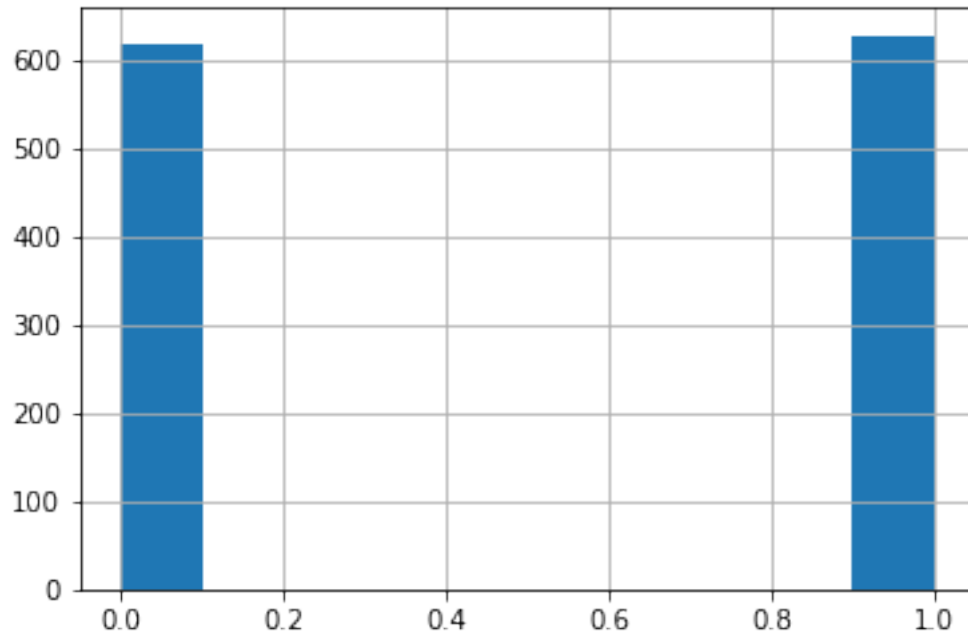
### 1.3 Question 3:

Apakah data loan (dataset 1) tersebut dikategorikan sebagai data yang akan mengalami masalah 'imbalanced class'?

**Answer:** Tidak

```
In [10]: df_train['not_paid'].hist()
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x22ccd141f28>
```



Menggunakan train dataset 1, buat sebuah model regresi logistik untuk memprediksi peluang default (not\_paid) nasabah dengan variabel purpose, int\_rate, installment, annual\_inc, verified, home\_ownership dan grdCtoA sebagai variabel prediktornya.

#### 1.4 Question 4:

Variabel manakah yang memiliki korelasi negatif terhadap kenaikan odds default (not\_paid=1)?

**Answer:**

```
In [11]: used_variables = ['purpose', 'int_rate', 'installment', 'annual_inc', 'verified', 'home_ownership', 'grdCtoA']
```

```
nasabah = df_train[used_variables]
```

```
In [12]: corr_matrix = nasabah.corr()
corr_matrix['not_paid'].sort_values(ascending=False)
```

```
Out[12]: not_paid      1.000000
int_rate      0.160727
installment    0.148875
verified       0.101174
annual_inc    -0.058832
grdCtoA       -0.153679
Name: not_paid, dtype: float64
```

## 2 Preprocessing

```
In [13]: nasabah.head()
```

```
Out[13]:
```

	purpose	int_rate	installment	annual_inc	verified	\
0	debt_consolidation	21.45	955.75	90000.0	1	
1	debt_consolidation	18.06	289.47	65000.0	1	
2	home_improvement	9.44	838.91	97400.0	1	
3	home_improvement	10.42	214.55	60000.0	0	
4	debt_consolidation	11.99	1024.52	150000.0	0	

	home_ownership	grdCtoA	not_paid
0	MORTGAGE	0	1
1	MORTGAGE	0	1
2	MORTGAGE	1	1
3	OWN	1	0
4	MORTGAGE	1	0

```
In [14]: # check missing values
nasabah.isnull().sum()
```

```
Out[14]: purpose          0
int_rate                0
installment            0
annual_inc             0
verified               0
home_ownership         0
grdCtoA                0
not_paid               0
dtype: int64
```

```
In [15]: nasabah_num = nasabah.drop(['purpose', 'home_ownership'], axis=1)
nasabah_num.head()
```

```
Out[15]:
```

	int_rate	installment	annual_inc	verified	grdCtoA	not_paid
0	21.45	955.75	90000.0	1	0	1
1	18.06	289.47	65000.0	1	0	1
2	9.44	838.91	97400.0	1	1	1
3	10.42	214.55	60000.0	0	1	0
4	11.99	1024.52	150000.0	0	1	0

```
In [16]: nasabah_cat = nasabah[['purpose', 'home_ownership']]
nasabah_cat.head()
```

```
Out[16]:
```

	purpose	home_ownership
0	debt_consolidation	MORTGAGE
1	debt_consolidation	MORTGAGE
2	home_improvement	MORTGAGE
3	home_improvement	OWN
4	debt_consolidation	MORTGAGE

```
In [17]: nasabah_cat['purpose'].value_counts()
```

```
Out[17]: debt_consolidation    808
         credit_card          247
         home_improvement      126
         major_purchase        40
         small_business        23
         Name: purpose, dtype: int64
```

```
In [18]: nasabah_cat['home_ownership'].value_counts()
```

```
Out[18]: MORTGAGE    633
         RENT        458
         OWN         153
         Name: home_ownership, dtype: int64
```

## 2.1 Change categorical using One Hot Encoding

```
In [19]: from future_encoders import OneHotEncoder
```

```
In [20]: purpose_cat = nasabah_cat[['purpose']]
         purpose_cat.head()
```

```
Out[20]:           purpose
0  debt_consolidation
1  debt_consolidation
2   home_improvement
3   home_improvement
4  debt_consolidation
```

```
In [21]: encoder = OneHotEncoder(sparse=False)
         purpose_cat_1hot = encoder.fit_transform(purpose_cat)
         purpose_cat_1hot
```

```
Out[21]: array([[0., 1., 0., 0., 0.],
                [0., 1., 0., 0., 0.],
                [0., 0., 1., 0., 0.],
                ...,
                [1., 0., 0., 0., 0.],
                [0., 1., 0., 0., 0.],
                [0., 1., 0., 0., 0.]])
```

```
In [22]: encoder.categories_
```

```
Out[22]: [array(['credit_card', 'debt_consolidation', 'home_improvement',
                'major_purchase', 'small_business'], dtype=object)]
```

## 2.2 Pipeline

```
In [23]: from sklearn.pipeline import Pipeline
         from sklearn.preprocessing import StandardScaler
         from sklearn.preprocessing import Imputer
         from future_encoders import OneHotEncoder, ColumnTransformer
```



```
In [24]: nasabah_num.head()
```

```
Out[24]:
```

	int_rate	installment	annual_inc	verified	grdCtoA	not_paid
0	21.45	955.75	90000.0	1	0	1
1	18.06	289.47	65000.0	1	0	1
2	9.44	838.91	97400.0	1	1	1
3	10.42	214.55	60000.0	0	1	0
4	11.99	1024.52	150000.0	0	1	0

```
In [25]: num_attribs = ['int_rate', 'installment', 'annual_inc']
bool_attribs = ['verified', 'grdCtoA']
cat_attribs = ['purpose', 'home_ownership']
```

```
full_pipeline = ColumnTransformer([
    ("num", StandardScaler(), num_attribs),
    ("bool", Imputer(strategy="median"), bool_attribs),
    ("cat", OneHotEncoder(sparse=False), cat_attribs)
])
```

```
nasabah_prepared = full_pipeline.fit_transform(nasabah)
```

```
In [26]: nasabah_prepared.shape
```

```
Out[26]: (1244, 13)
```

```
In [27]: X_train = nasabah_prepared
y_train = nasabah['not_paid']
```

### 3 Logistic Regression Model

```
In [28]: from sklearn.linear_model import LogisticRegression
```

```
log_reg = LogisticRegression(random_state=42)
log_reg.fit(X_train, y_train)
```

```
Out[28]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=42, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)
```

### 4 Fine-tune Model

```
In [29]: from sklearn.model_selection import GridSearchCV
```

```
In [51]: # regularization penalty space
penalty = ['l1', 'l2']
```

```
# regularization hyperparameter space
```

```

C = np.logspace(0, 1, 10)

# create hyperparameter options
hyperparameters = dict(C=C, penalty=penalty)

In [52]: grid_search = GridSearchCV(log_reg, hyperparameters, cv=5, scoring='neg_mean_squared_

grid_search.fit(X_train, y_train)

Out[52]: GridSearchCV(cv=5, error_score='raise',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=42, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False),
    fit_params=None, iid=True, n_jobs=1,
    param_grid={'C': array([ 1.        ,  1.29155,  1.6681 ,  2.15443,  2.78256,  3.59381,
    4.64159,  5.99484,  7.74264, 10.        ]), 'penalty': ['l1', 'l2']},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring='neg_mean_squared_error', verbose=0)

In [53]: grid_search.best_params_

Out[53]: {'C': 1.0, 'penalty': 'l1'}

In [54]: grid_search.best_estimator_

Out[54]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l1', random_state=42, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)

In [55]: cvres = grid_search.cv_results_
    for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
        print(np.sqrt(-mean_score), params)

0.6452896356105269 {'C': 1.0, 'penalty': 'l1'}
0.6502534751876141 {'C': 1.0, 'penalty': 'l2'}
0.6477763100870298 {'C': 1.2915496650148839, 'penalty': 'l1'}
0.6508712933556599 {'C': 1.2915496650148839, 'penalty': 'l2'}
0.6490160744941288 {'C': 1.6681005372000588, 'penalty': 'l1'}
0.6508712933556599 {'C': 1.6681005372000588, 'penalty': 'l2'}
0.6490160744941288 {'C': 2.154434690031884, 'penalty': 'l1'}
0.6502534751876141 {'C': 2.154434690031884, 'penalty': 'l2'}
0.6490160744941288 {'C': 2.7825594022071245, 'penalty': 'l1'}
0.6502534751876141 {'C': 2.7825594022071245, 'penalty': 'l2'}
0.6502534751876141 {'C': 3.5938136638046276, 'penalty': 'l1'}
0.6502534751876141 {'C': 3.5938136638046276, 'penalty': 'l2'}
0.6508712933556599 {'C': 4.641588833612778, 'penalty': 'l1'}
0.6502534751876141 {'C': 4.641588833612778, 'penalty': 'l2'}

```

```

0.6502534751876141 {'C': 5.994842503189409, 'penalty': 'l1'}
0.6502534751876141 {'C': 5.994842503189409, 'penalty': 'l2'}
0.6502534751876141 {'C': 7.742636826811269, 'penalty': 'l1'}
0.6502534751876141 {'C': 7.742636826811269, 'penalty': 'l2'}
0.6502534751876141 {'C': 10.0, 'penalty': 'l1'}
0.6502534751876141 {'C': 10.0, 'penalty': 'l2'}

```

## 5 Test the model

```
In [35]: df_test = pd.read_csv('dataset1/test1.csv')
```

```
In [36]: df_test.head()
```

```

Out[36]:   Unnamed: 0  initial_list_status      purpose  int_rate  installment  \
0           6                    w  debt_consolidation    10.91         130.79
1           8                    w      credit_card    10.91         915.50
2           9                    w   home_improvement    17.09         713.96
3          10                    w  debt_consolidation    18.06         408.73
4          26                    w  debt_consolidation    18.06         578.93

   annual_inc  dti  verification_status  grade  revol_bal  inq_last_12m  \
0   49000.0    5.12      Not Verified      B      2016         5
1   95000.0   33.11      Not Verified      B     27588         1
2  150000.0   14.26   Source Verified      D     27024         8
3   85000.0   17.66      Verified      D     11719         1
4   40000.0   25.32   Source Verified      D     15264         2

   delinq_2yrs  home_ownership  not_paid  log_inc  verified  grdCtoA
0           0      MORTGAGE          1  10.799576          0          1
1           0          RENT          1  11.461632          0          1
2           0      MORTGAGE          1  11.918391          1          0
3           0          RENT          0  11.350407          1          0
4           0          RENT          1  10.596635          1          0

```

```
In [37]: df_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 312 entries, 0 to 311
Data columns (total 17 columns):
Unnamed: 0      312 non-null int64
initial_list_status  312 non-null object
purpose         312 non-null object
int_rate        312 non-null float64
installment     312 non-null float64
annual_inc      312 non-null float64
dti             312 non-null float64
verification_status  312 non-null object

```

```

grade                312 non-null object
revol_bal             312 non-null int64
inq_last_12m         312 non-null int64
delinq_2yrs          312 non-null int64
home_ownership        312 non-null object
not_paid             312 non-null int64
log_inc              312 non-null float64
verified             312 non-null int64
grdCtoA              312 non-null int64
dtypes: float64(5), int64(7), object(5)
memory usage: 41.5+ KB

```

```

In [38]: nasabah_test = df_test[used_variables]
        nasabah_test.head()

```

```

Out[38]:
   purpose  int_rate  installment  annual_inc  verified \
0  debt_consolidation    10.91      130.79    49000.0        0
1    credit_card        10.91      915.50    95000.0        0
2  home_improvement    17.09      713.96   150000.0        1
3  debt_consolidation    18.06      408.73    85000.0        1
4  debt_consolidation    18.06      578.93    40000.0        1

   home_ownership  grdCtoA  not_paid
0      MORTGAGE        1         1
1         RENT        1         1
2      MORTGAGE        0         1
3         RENT        0         0
4         RENT        0         1

```

```

In [39]: X_test = full_pipeline.fit_transform(nasabah_test)

```

```

In [40]: X_test.shape

```

```

Out[40]: (312, 13)

```

## 6 Accuracy Score

```

In [41]: from sklearn.metrics import accuracy_score

```

```

In [42]: y_true = nasabah_test['not_paid']

```

```

In [43]: final_model = grid_search.best_estimator_
        predict = final_model.predict(X_test)

```

```

In [44]: accuracy_score(y_true, predict)

```

```

Out[44]: 0.5993589743589743

```

```
In [45]: from sklearn.metrics import confusion_matrix

        tn, fp, fn, tp = confusion_matrix(y_true, predict).ravel()

        print(tn, fp, fn, tp)

91 70 55 96
```

## 6.1 Question 5:

Jika pada model regresi logistik diperoleh koefisien dari variabel gradeCtoA adalah -0.3298. Dengan mengasumsikan variabel lain konstan, berapa odds default (not\_paid=1) untuk nasabah yang memiliki grad A-C (gradeCtoA=1) dibandingkan dengan nasabah yang memiliki grade loan lain? bulatkan hasil anda hingga 3 angka dibelakang koma (contoh : 4.323 atau 16.423)

**Answer:**

```
In [46]: pass
```

## 6.2 Question 6

Diberikan cross tabulasi hasil perbandingan nilai aktual dan prediksi menggunakan model regresi logistik berikut :

		actual	
		0	1
predicted	0	93	54
	1	68	97

Berapa nilai recall/sensitivity berdasarkan confusion matrix diatas?

**Answer:**  $97/(97+54) = 0.6423841059602649$

```
In [56]: recall = tp/(tp+fn)
        recall
```

```
Out [56]: 0.6357615894039735
```